

```
In [1]: %matplotlib inline
```

1 Universal adversarial training & defence

A modification of the model proposed in [Universal adversarial training](#) by Shafahi et al., where we aim to simultaneously find the optimal attack δ and defence for a given network

Author(s): [Martin Benning](#), Alex Wendland

Date: 11.04.2019

Last modified: 11.04.2019

```
In [2]: from __future__ import print_function
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.autograd import Variable
from torchvision import datasets, transforms
import numpy as np
import matplotlib.pyplot as plt

pretrained_model = "../data/lenet_mnist_model.pth"

In [3]: # LeNet Model definition
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.conv2_drop = nn.Dropout2d()
        self.fc1 = nn.Linear(320, 50)
        self.fc2 = nn.Linear(50, 10)

    def forward(self, x):
        x = F.relu(F.max_pool2d(self.conv1(x), 2))
        x = F.relu(F.max_pool2d(self.conv2_drop(self.conv2(x)), 2))
        x = x.view(-1, 320)
        x = F.relu(self.fc1(x))
        x = F.dropout(x, training=self.training)
        x = self.fc2(x)
        return F.log_softmax(x, dim=1)

# MNIST Train & test dataset and dataloader declaration

train_loader = torch.utils.data.DataLoader(
    datasets.MNIST('../data/MNIST', train=True, download=True,
        transform=transforms.Compose([
            transforms.ToTensor(),
        ])),
```

```

        batch_size=6000, shuffle=True)

test_loader = torch.utils.data.DataLoader(
    datasets.MNIST('../data/MNIST', train=False, download=True,
        transform=transforms.Compose([
            transforms.ToTensor(),
        ])),
    batch_size=1, shuffle=True)

# Initialize the network
model = Net()

# Load the pretrained model
# model.load_state_dict(torch.load(pretrained_model, map_location='cpu'))

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to
../data/MNIST/MNIST/raw/train-images-idx3-ubyte.gz

In [4]: perturbation = Variable(torch.zeros(train_loader.dataset.train_data.size(1), \
        train_loader.dataset.train_data.size(2)), \
        requires_grad=True)

no_of_epochs = 150

stepsize = 0.1
optimiser = optim.SGD(model.parameters(), lr=stepsize, momentum=0.7)
epsilon = 3

for epoch in range(no_of_epochs):

    for data, target in train_loader:

        output_1 = model(data)
        output_2 = model(data + perturbation.repeat(train_loader.batch_size, \
            1, 1, 1))

        loss_1 = F.nll_loss(output_1, target)
        loss_2 = F.nll_loss(output_2, target)
        loss = 1/2 * (loss_1 + loss_2)
        model.zero_grad()
        loss.backward()

        optimiser.step()

    perturbation.data = perturbation.data + stepsize * perturbation.grad
    perturbation.data = perturbation.data \
        / torch.norm(perturbation.data.view(-1, 784)) \

```

```

        * epsilon

    print('Iteration [%d/%d], Loss: %.4f'
          %(epoch + 1, no_of_epochs, loss.item()))

    print('Iteration [%d/%d] completed, Loss: %.4f'
          %(epoch + 1, no_of_epochs, loss.item()))

Iteration [150/150] completed, Loss: 0.1512


```

```

In [5]: plt.imshow(perturbation.data.numpy())
        plt.savefig('../Figures/attack_defence_perturbation.png')
        print(torch.norm(perturbation.data.view(-1, 784)))

tensor(3.)

```



```

In [9]: # Set the model in evaluation mode. In this case this is for the Dropout layers
        model.eval()

        dataiter = iter(test_loader)
        images, labels = dataiter.next()
        plt.imshow(images[0][0].data.numpy())
        plt.savefig('../Figures/attack_defence_image.png')
        plt.imshow(images[0][0].data.numpy() + perturbation.data.numpy())


```

```

plt.savefig('../Figures/attack_defence_perturbed_image.png')
output1 = model(images)
pred1 = output1.max(1, keepdim=True)[1]
output2 = model(images + perturbation)
pred2 = output2.max(1, keepdim=True)[1]
print(labels, pred1, pred2)

tensor([0]) tensor([[0]]) tensor([[0]])

```



```

In [10]: def test_accuracy(model, test_loader, perturbation):

    # Accuracy counter
    correct = 0
    adv_examples = []

    # Loop over all examples in test set
    for data, target in test_loader:

        # Forward pass the data through the model
        perturbed_data = data + perturbation
        output = model(perturbed_data)

        # Check for success
        final_pred = output.max(1, keepdim=True)[1]
        # get the index of the max log-probability

```

```

    if final_pred.item() == target.item():
        correct += 1
    else:
        # Save some adv examples for visualisation later
        if len(adv_examples) < 5:
            adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
            adv_examples.append( (final_pred.item(), adv_ex) )

        # Calculate final accuracy for this epsilon
        final_acc = correct/float(len(test_loader))
        print("Test Accuracy = {} / {} = {}".format(correct, \
            len(test_loader), final_acc))

        # Return the accuracy and an adversarial example
        return final_acc, adv_examples

```

```

In [11]: acc_1, ex_1 = test_accuracy(model, test_loader, torch.zeros(28, 28))
         acc_2, ex_2 = test_accuracy(model, test_loader, perturbation)

```

Test Accuracy = 9849 / 10000 = 0.9849

Test Accuracy = 9798 / 10000 = 0.9798

```

In [12]: # Initialize the network
         model_pretrained = Net()

         # Load the pretrained model
         model_pretrained.load_state_dict(torch.load(pretrained_model, \
             map_location='cpu'))

         acc_3, ex_3 = test_accuracy(model_pretrained, test_loader, perturbation)

```

Test Accuracy = 7491 / 10000 = 0.7491